

## Project 5 - OpenCL Array Multiply, Multiply-Add, and Multiply-Reduce

### Multiply and Multiply-Add

1. What machine you ran this on  
Rabbit
2. Show the tables and graphs

#### Multiply

# of Elements	Local Size	# of Work Groups	GigaMultsPerSecond
1024	8	128	0.0284966
2048	8	256	0.0595006
4096	8	512	0.119169
8192	8	1024	0.239063
16384	8	2048	0.231696
32768	8	4096	0.904297
65536	8	8192	0.783896
131072	8	16384	3.67826
262144	8	32768	1.40423
524288	8	65536	2.31552
1048576	8	131072	4.88003
2097152	8	262144	9.02098
4194304	8	524288	16.4147
8388608	8	1048576	28.9842
1024	16	64	0.0302762
2048	16	128	0.0330452

Project 5 - OpenCL Array Multiply, Multiply-Add, and Multiply-Reduce

4096	16	256	0.122283
8192	16	512	0.114072
16384	16	1024	0.233752
32768	16	2048	0.427826
65536	16	4096	1.65449
131072	16	8192	3.87195
262144	16	16384	0.696689
524288	16	32768	1.73423
1048576	16	65536	5.43194
2097152	16	131072	6.63195
4194304	16	262144	15.2991
8388608	16	524288	28.5838
1024	32	32	0.0147372
2048	32	64	0.0248978
4096	32	128	0.114312
8192	32	256	0.240239
16384	32	512	0.467653
32768	32	1024	0.945156
65536	32	2048	1.97743
131072	32	4096	1.78442
262144	32	8192	0.796306
524288	32	16384	2.63189
1048576	32	32768	5.50223
2097152	32	65536	8.84905

Project 5 - OpenCL Array Multiply, Multiply-Add, and Multiply-Reduce

4194304	32	131072	17.503
8388608	32	262144	27.8406
1024	64	16	0.0290922
2048	64	32	0.0588352
4096	64	64	0.0572409
8192	64	128	0.233665
16384	64	256	0.473647
32768	64	512	0.452765
65536	64	1024	1.8754
131072	64	2048	3.78123
262144	64	4096	1.33288
524288	64	8192	1.90687
1048576	64	16384	5.35592
2097152	64	32768	9.95156
4194304	64	65536	16.0752
8388608	64	131072	28.2262
1024	128	8	0.0158044
2048	128	16	0.0583977
4096	128	32	0.117281
8192	128	64	0.232972
16384	128	128	0.233665
32768	128	256	0.882787
65536	128	512	1.91667
131072	128	1024	3.82418

Project 5 - OpenCL Array Multiply, Multiply-Add, and Multiply-Reduce

262144	128	2048	1.36941
524288	128	4096	2.36188
1048576	128	8192	5.32718
2097152	128	16384	9.65683
4194304	128	32768	17.0868
8388608	128	65536	28.3201
1024	256	4	0.0299741
2048	256	8	0.0594395
4096	256	16	0.117256
8192	256	32	0.240646
16384	256	64	0.23082
32768	256	128	0.607234
65536	256	256	1.85122
131072	256	512	3.33849
262144	256	1024	1.33655
524288	256	2048	2.3961
1048576	256	4096	5.56027
2097152	256	8192	9.55253
4194304	256	16384	16.3068
8388608	256	32768	28.6226
1024	512	2	0.0286719
2048	512	4	0.0574187
4096	512	8	0.120839
8192	512	16	0.115805

Project 5 - OpenCL Array Multiply, Multiply-Add, and Multiply-Reduce

16384	512	32	0.452963
32768	512	64	0.72369
65536	512	128	0.866546
131072	512	256	3.7562
262144	512	512	1.32606
524288	512	1024	2.34756
1048576	512	2048	4.46959
2097152	512	4096	10.8745
4194304	512	8192	13.907
8388608	512	16384	25.9464

Multiply-Add

# of Elements	Local Size	# of Work Groups	GigaMultsPerSecond
1024	8	128	0.0299709
2048	8	256	0.05881
4096	8	512	0.0590564
8192	8	1024	0.112126
16384	8	2048	0.476237
32768	8	4096	0.455614
65536	8	8192	1.91251
131072	8	16384	3.78408
262144	8	32768	1.0397
524288	8	65536	2.7177

Project 5 - OpenCL Array Multiply, Multiply-Add, and Multiply-Reduce

1048576	8	131072	4.13129
2097152	8	262144	10.5679
4194304	8	524288	13.7916
8388608	8	1048576	27.9647
1024	16	64	0.0290953
2048	16	128	0.0578995
4096	16	256	0.123333
8192	16	512	0.238519
16384	16	1024	0.46305
32768	16	2048	0.919036
65536	16	4096	0.927613
131072	16	8192	2.36351
262144	16	16384	0.94113
524288	16	32768	2.7624
1048576	16	65536	5.61277
2097152	16	131072	9.68191
4194304	16	262144	12.1143
8388608	16	524288	27.4854
1024	32	32	0.0302963
2048	32	64	0.0604559
4096	32	128	0.116166
8192	32	256	0.23754
16384	32	512	0.23296
32768	32	1024	0.936801

Project 5 - OpenCL Array Multiply, Multiply-Add, and Multiply-Reduce

65536	32	2048	1.87022
131072	32	4096	3.7242
262144	32	8192	1.49066
524288	32	16384	2.73224
1048576	32	32768	5.71071
2097152	32	65536	7.1273
4194304	32	131072	15.8852
8388608	32	262144	24.3155
1024	64	16	0.0291385
2048	64	32	0.0569607
4096	64	64	0.120972
8192	64	128	0.116246
16384	64	256	0.480346
32768	64	512	0.932629
65536	64	1024	1.91971
131072	64	2048	2.74899
262144	64	4096	1.5012
524288	64	8192	2.75632
1048576	64	16384	5.62792
2097152	64	32768	10.8939
4194304	64	65536	15.0494
8388608	64	131072	26.8743
1024	128	8	0.0299872
2048	128	16	0.053209

Project 5 - OpenCL Array Multiply, Multiply-Add, and Multiply-Reduce


4096	128	32	0.0575164
8192	128	64	0.232197
16384	128	128	0.479874
32768	128	256	0.946987
65536	128	512	1.91064
131072	128	1024	3.87067
262144	128	2048	1.47747
524288	128	4096	1.94436
1048576	128	8192	5.3228
2097152	128	16384	10.8088
4194304	128	32768	16.5574
8388608	128	65536	27.7519
1024	256	4	0.0145801
2048	256	8	0.0579911
4096	256	16	0.0583249
8192	256	32	0.239336
16384	256	64	0.443732
32768	256	128	0.967667
65536	256	256	1.88061
131072	256	512	3.88692
262144	256	1024	1.45989
524288	256	2048	2.22684
1048576	256	4096	3.98984
2097152	256	8192	10.9621



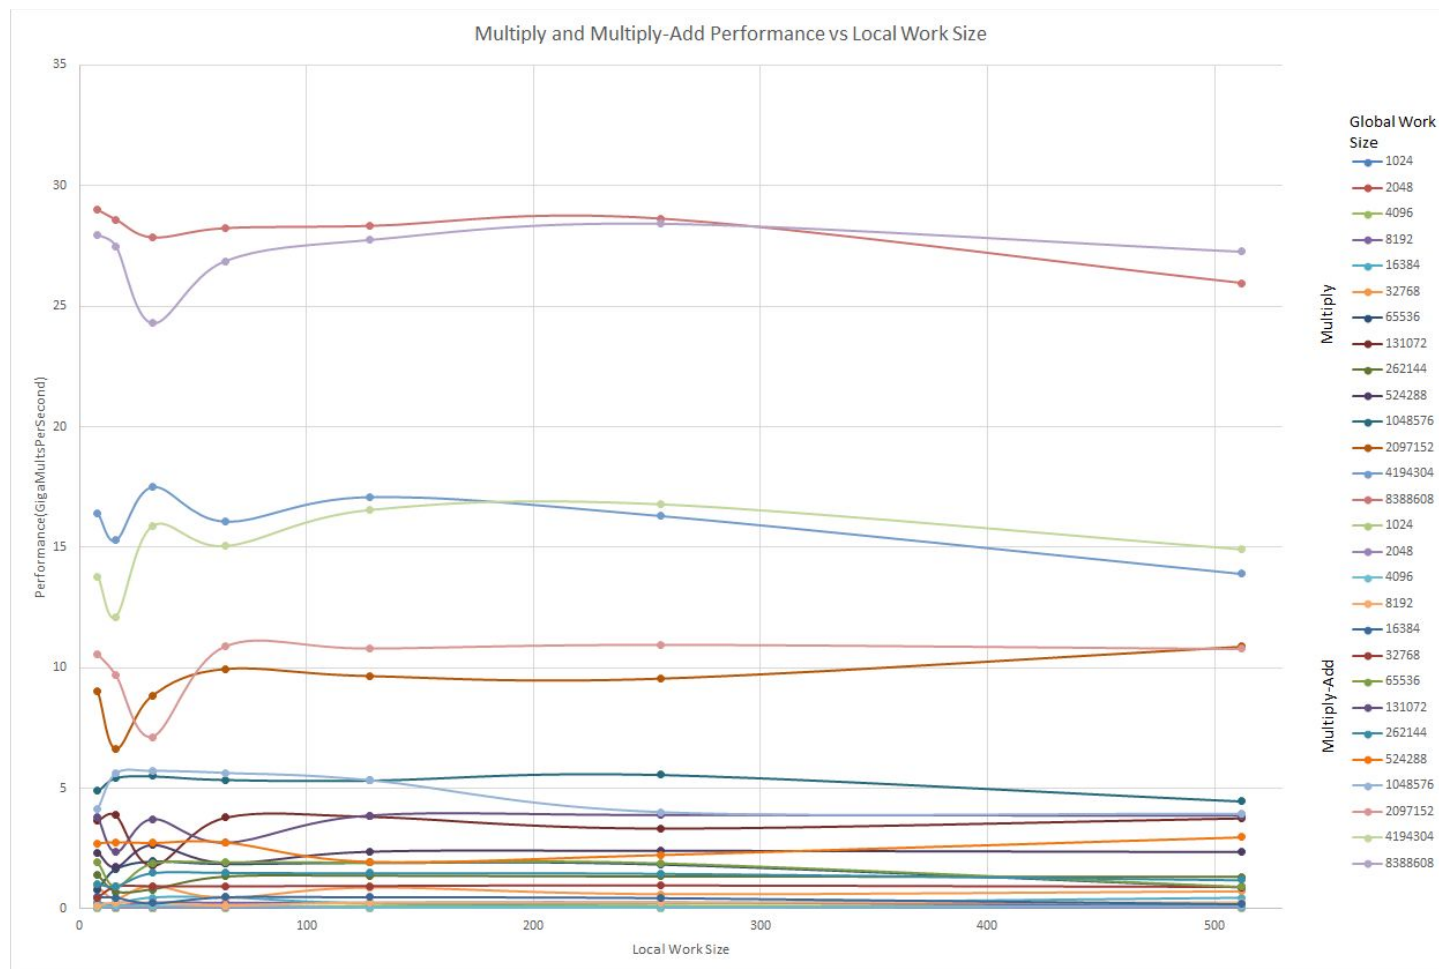
Project 5 - OpenCL Array Multiply, Multiply-Add, and Multiply-Reduce

4194304	256	16384	16.7891
8388608	256	32768	28.4169
1024	512	2	0.0146488
2048	512	4	0.0599026
4096	512	8	0.117269
8192	512	16	0.240488
16384	512	32	0.198428
32768	512	64	0.904762
65536	512	128	0.924032
131072	512	256	3.86069
262144	512	512	1.19362
524288	512	1024	2.96442
1048576	512	2048	3.92127
2097152	512	4096	10.7882
4194304	512	8192	14.9133
8388608	512	16384	27.2613

## Project 5 - OpenCL Array Multiply, Multiply-Add, and Multiply-Reduce



## Project 5 - OpenCL Array Multiply, Multiply-Add, and Multiply-Reduce



### 3. What patterns are you seeing in the performance curves?

In general, it appears that performance is directly dependent on global work size. It seems that the greater the global work size, the better the performance, with some minor variations due to local work size. Performance between multiply and multiply-add was close. There doesn't seem to be a clear trend indicating if multiply or multiply-add is more performant. Local work size doesn't seem to make a significant impact either. It seems as though smaller local work sizes produce some squirrely behavior, but performance levels out at around 64 for local work size. Performance starts to marginally decline at around a local work size of 256.

### 4. Why do you think the patterns look this way?

I think the biggest, most conclusive pattern of note is that the performance is dependent on global work size. This makes sense because we are using the GPU to perform the same calculation over and over again. This is what GPU's are built for. The more data we give the GPU, the faster it seems to tear through the data. I suspect that this would only go to a certain point due to there being a limited number of cores on a GPU. I don't think I reached that point because rabbit has impressive specifications.

### 5. What is the performance difference between doing a Multiply and doing a Multiply-Add?

Performance difference between multiply and multiply-add is inconclusive at best for my results. There are times when multiply is more performant and there seem to be equal times where multiply-add is

Project 5 - OpenCL Array Multiply, Multiply-Add, and Multiply-Reduce

more performant. I think due to the nature of GPU being made for these kinds of operations, it rips through the operations on the data no matter if there is an extra step or not.

**6. What does that mean for the proper use of GPU parallel computing?**

I suspect that I should have gotten that multiply is more performant than multiply-add because multiply-add has to store intermediate values, which storage is not the GPU's forte. It makes sense in theory that a GPU would perform a straight multiply faster than a multiply-add because of this. My results didn't seem to prove the theory.

## Multiply-Reduction

**1. Show this table and graph**

# of Elements	Local Size	# of Work Groups	GigaMultsPerSecond
1024	8	128	0.00316574
2048	8	256	0.0112834
4096	8	512	0.0200224
8192	8	1024	0.0410028
16384	8	2048	0.0854786
32768	8	4096	0.149064
65536	8	8192	0.247674
131072	8	16384	0.49418
262144	8	32768	0.599329
524288	8	65536	0.846765
1048576	8	131072	0.692291
2097152	8	262144	0.669207
4194304	8	524288	1.18786
8388608	8	1048576	1.22161
1024	16	64	0.00574542
2048	16	128	0.0111002

Project 5 - OpenCL Array Multiply, Multiply-Add, and Multiply-Reduce

4096	16	256	0.0225206
8192	16	512	0.0405133
16384	16	1024	0.0865578
32768	16	2048	0.151582
65536	16	4096	0.251285
131072	16	8192	0.438104
262144	16	16384	0.889679
524288	16	32768	0.911911
1048576	16	65536	1.39817
2097152	16	131072	0.791572
4194304	16	262144	1.93612
8388608	16	524288	2.0737
1024	32	32	0.00297308
2048	32	64	0.0113291
4096	32	128	0.0229032
8192	32	256	0.0444314
16384	32	512	0.0900612
32768	32	1024	0.165751
65536	32	2048	0.247875
131072	32	4096	0.468691
262144	32	8192	0.887458
524288	32	16384	1.44201
1048576	32	32768	1.90108
2097152	32	65536	1.01379

Project 5 - OpenCL Array Multiply, Multiply-Add, and Multiply-Reduce

4194304	32	131072	3.20771
8388608	32	262144	3.34469
1024	64	16	0.003963
2048	64	32	0.00900707
4096	64	64	0.0211469
8192	64	128	0.0441855
16384	64	256	0.0833255
32768	64	512	0.138783
65536	64	1024	0.319762
131072	64	2048	0.0314633
262144	64	4096	1.15912
524288	64	8192	1.91301
1048576	64	16384	2.94328
2097152	64	32768	1.35219
4194304	64	65536	4.61609
8388608	64	131072	5.4087
1024	128	8	0.00362072
2048	128	16	0.0103215
4096	128	32	0.0210937
8192	128	64	0.0417827
16384	128	128	0.0688453
32768	128	256	0.159729
65536	128	512	0.337532
131072	128	1024	0.466584

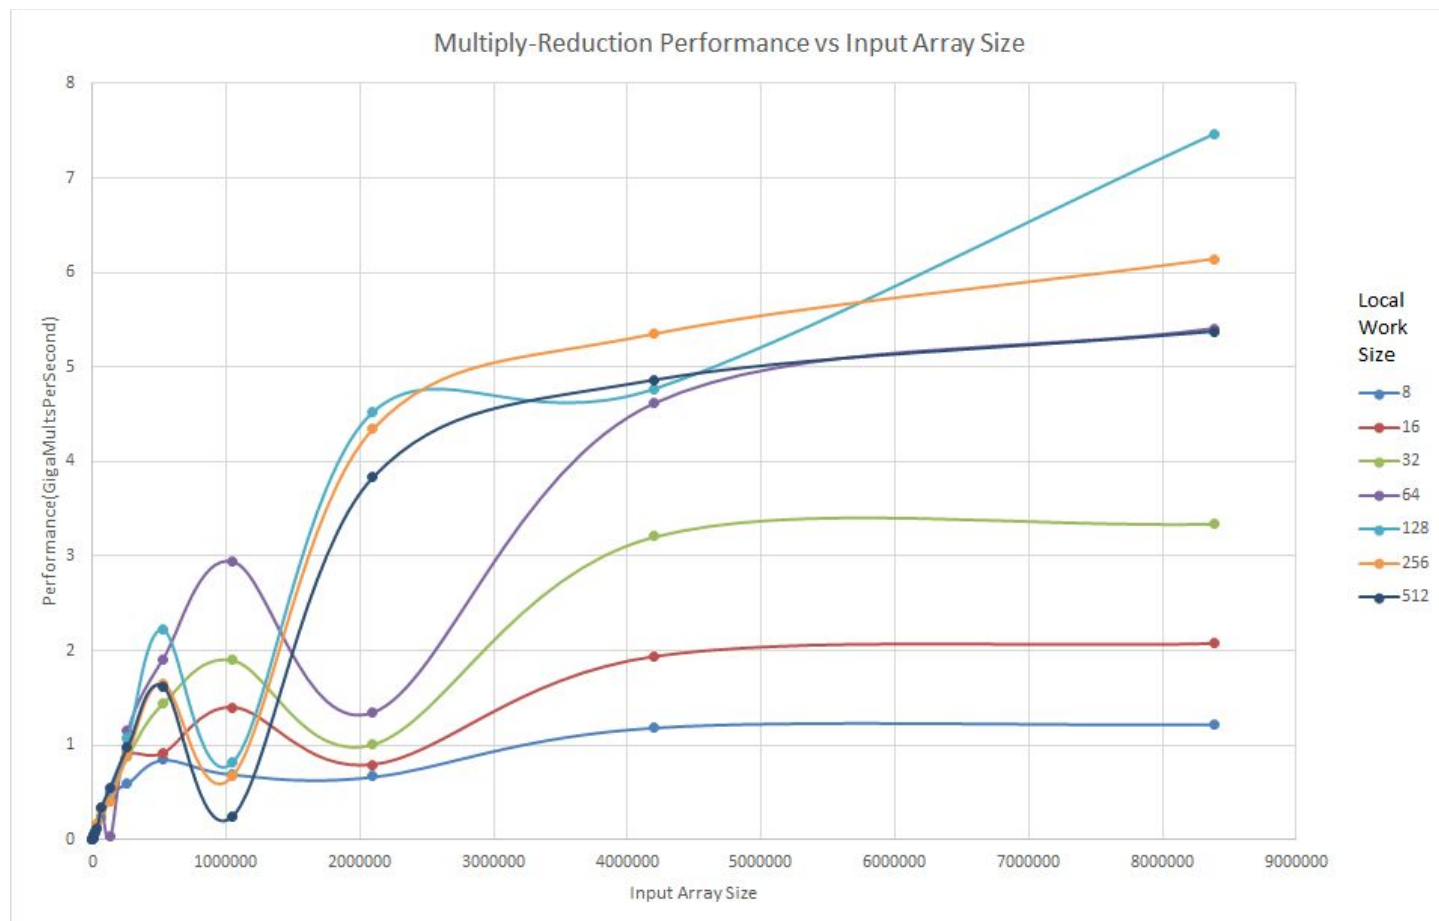
Project 5 - OpenCL Array Multiply, Multiply-Add, and Multiply-Reduce

262144	128	2048	1.07399
524288	128	4096	2.22573
1048576	128	8192	0.827015
2097152	128	16384	4.52978
4194304	128	32768	4.76915
8388608	128	65536	7.46744
1024	256	4	0.0056917
2048	256	8	0.0107596
4096	256	16	0.020194
8192	256	32	0.0453402
16384	256	64	0.0828905
32768	256	128	0.163203
65536	256	256	0.326561
131072	256	512	0.399496
262144	256	1024	0.899293
524288	256	2048	1.64783
1048576	256	4096	0.673281
2097152	256	8192	4.34465
4194304	256	16384	5.34891
8388608	256	32768	6.14375
1024	512	2	0.00557488
2048	512	4	0.0111204
4096	512	8	0.01819
8192	512	16	0.0455445

Project 5 - OpenCL Array Multiply, Multiply-Add, and Multiply-Reduce

16384	512	32	0.0903117
32768	512	64	0.120872
65536	512	128	0.342667
131072	512	256	0.550138
262144	512	512	0.982056
524288	512	1024	1.62062
1048576	512	2048	0.249313
2097152	512	4096	3.83237
4194304	512	8192	4.85759
8388608	512	16384	5.37222





## 2. What pattern are you seeing in this performance curve?

Overall, performance appears to trend up the larger the input array size gets. Interestingly, the local work size of 128 is more performant than any other for very large input array sizes. A local size of 256 comes mostly dominates in performance for all most input array sizes. Smaller local work sizes do not appear to perform as well as their counterparts. As usual, there is some squirrely behavior for smaller input array sizes. In general, performance is *much slower* than multiply and multiply-add.

## 3. Why do you think the pattern looks this way?

Once again, because we are using the GPU which excels at performing the same calculation on tons of data, we see that performance gets better with larger input array sizes. I think there is a sweet spot when it comes to local work size that is reliant on the GPU's number of cores. If the local work size is too high or low, performance will suffer. The GPU's number of cores are finite, and thus has an optimum for local work size in order to process massive amounts of data.

## 4. What does that mean for the proper use of GPU parallel computing?

For simpler operations, GPU parallel computing is unparalleled(see what I did there?) in terms of speeding up performance. It makes a tremendous impact on performance. Adding the aspect of a reduction in this example slowed performance significantly because GPU's are not intended to store

Project 5 - OpenCL Array Multiply, Multiply-Add, and Multiply-Reduce

data as CPU's are. It seems that the simpler the operation being performed, and the more data, the better the GPU will process the data.

**Summary:**

I really enjoyed this project. I learned so much about GPU's. I honestly had no clue how they operate, but just knew that they made my video games look cool. It makes a lot of sense that they are best to perform quick calculations because anything visual relies on that speed to look better. I never would've thought that's what happens behind the scenes. It's interesting to think that the GPU can be used for things that are other than visual in nature, such as machine learning. It also got me thinking about bitcoin mining, which I've heard utilizes GPU's quite a bit.